

# CS 480/680 Winter 2024:

## Lecture Notes

<b>1</b>	<b>Classic Machine Learning</b>	<b>2</b>
1.1	Introduction . . . . .	2
1.2	Perceptron . . . . .	4

Lecture notes taken, unless otherwise specified, by myself during section 002 of the Winter 2024 offering of CS 480/680, taught by Hongyang Zheng.

<b>Lectures</b>	<b>Lecture 2</b>	<b>Jan 11</b>	<b>2</b>
Lecture 1	Jan 9		2

# Chapter 1

## Classic Machine Learning

### 1.1 Introduction

There have been three historical AI booms:

*Lecture 1*  
*Jan 9*

1. 1950s–1970s: search-based algorithms (e.g., chess), failed when they realized AI is actually a hard problem
2. 1980s–1990s: expert systems
3. 2012 – present: deep learning

Machine learning is the subset of AI where a program can learn from experience.

Major learning paradigms of machine learning:

- Supervised learning: teacher/human labels answers (e.g., classification, ranking, etc.)
- Unsupervised learning: without labels (e.g., clustering, representation, generation, etc.)
- Reinforcement learning: rewards given for actions (e.g., gaming, pricing, etc.)
- Others: semi-supervised, active learning, etc.

Active focuses in machine learning research:

- Representation: improving the encoding of data into a space
- Generalization: improving the use of the model on new distributions
- Interpretation: understanding how deep learning actually works
- Complexity: improving time/space requirements
- Efficiency: reducing the amount of samples required
- Privacy: respecting legal/ethical concerns of data sourcing
- Robustness: gracefully failing under errors or malicious attack
- Applications

A machine learning algorithm has three phases: training, prediction, and evaluation.

*Lecture 2*  
*Jan 11*

**Definition 1.1.1** (dataset)

A dataset consists of a list of features  $\mathbf{x}_1, \dots, \mathbf{x}_n, \mathbf{x}'_1, \dots, \mathbf{x}'_m \in \mathbb{R}^d$  which are  $d$ -dimensional vectors and a label vector  $\mathbf{y}^\top \in \mathbb{R}^n$ .

Each training sample  $\mathbf{x}_i$  is associated with a label  $y_i$ . A test sample  $\mathbf{x}'_i$  may or may not be labelled.

**Example 1.1.2** (email filtering). Suppose we have a list  $D$  of  $d$  English words.

Define the training set  $X = [\mathbf{x}_1, \dots, \mathbf{x}_n] \in \mathbb{R}^{d \times n}$  and  $\mathbf{y} = [y_1, \dots, y_n] \in \{\pm 1\}^n$  such that  $\mathbf{x}_{ij} = 1$  if the word  $j \in D$  appears in email  $i$  (this is the bag-of-words representation):


	$\mathbf{x}_1$	$\mathbf{x}_2$	$\mathbf{x}_3$	$\mathbf{x}_4$	$\mathbf{x}_5$	$\mathbf{x}_6$	$\mathbf{x}'$
and	1	0	0	1	1	1	1
viagra	1	0	1	0	0	0	1
the	0	1	1	0	1	1	0
of	1	1	0	1	0	1	0
nigeria	1	0	0	0	1	0	0
$y$	+	-	+	-	+	-	?

Then, given a new email  $\mathbf{x}'_1$ , we must determine if it is spam or not.

**Example 1.1.3** (OR dataset). We want to train the OR function:

	$\mathbf{x}_1$	$\mathbf{x}_2$	$\mathbf{x}_3$	$\mathbf{x}_4$
	0	1	0	1
	0	0	1	1
$y$	-	+	+	+

This can be represented graphically by finding a line dividing the points:



## 1.2 Perceptron

### Definition 1.2.1

The inner product of vectors  $\langle \mathbf{a}, \mathbf{b} \rangle$  is the sum of the element-wise product  $\sum_j a_j b_j$ .

A linear function is a function  $f : \mathbb{R}^d \rightarrow \mathbb{R}^d$  such that for all  $\alpha, \beta \in \mathbb{R}$ ,  $\mathbf{x}, \mathbf{z} \in \mathbb{R}^d$ ,  $f(\alpha\mathbf{x} + \beta\mathbf{z}) = \alpha f(\mathbf{x}) + \beta f(\mathbf{z})$ .

### Theorem 1.2.2 (linear duality)

A function is linear if and only if there exists  $\mathbf{w} \in \mathbb{R}^d$  such that  $f(\mathbf{x}) = \langle \mathbf{x}, \mathbf{w} \rangle$ .

*Proof.* ( $\Rightarrow$ ) Suppose  $f$  is linear. Let  $\mathbf{w} := [f(\mathbf{e}_1), \dots, f(\mathbf{e}_d)]$  where  $\mathbf{e}_i$  are coordinate vectors. Then:

$$\begin{aligned} f(\mathbf{x}) &= f(x_1\mathbf{e}_1 + \dots + x_d\mathbf{e}_d) \\ &= x_1f(\mathbf{e}_1) + \dots + x_df(\mathbf{e}_d) \\ &= \langle \mathbf{x}, \mathbf{w} \rangle \end{aligned}$$

by linearity of  $f$ .

( $\Leftarrow$ ) Suppose there exists  $\mathbf{w}$  such that  $f(\mathbf{x}) = \langle \mathbf{x}, \mathbf{w} \rangle$ . Then:

$$\begin{aligned} f(\alpha\mathbf{x} + \beta\mathbf{z}) &= \langle \alpha\mathbf{x} + \beta\mathbf{z}, \mathbf{w} \rangle \\ &= \alpha \langle \mathbf{x}, \mathbf{w} \rangle + \beta \langle \mathbf{z}, \mathbf{w} \rangle \\ &= \alpha f(\mathbf{x}) + \beta f(\mathbf{z}) \end{aligned}$$

since inner products are linear in the first argument. □

### Definition 1.2.3 (affine function)

A function  $f(\mathbf{x})$  where there exist  $\mathbf{w} \in \mathbb{R}^d$  and bias  $b \in \mathbb{R}$  such that  $f(\mathbf{x}) = \langle \mathbf{x}, \mathbf{w} \rangle + b$ .

### Definition 1.2.4 (sign function)

$$\text{sgn}(t) = \begin{cases} +1 & t > 0 \\ -1 & t \leq 0 \end{cases}$$


It does not matter what  $\text{sgn}(0)$  is defined as.

### Definition 1.2.5 (linear classifier)

$$\hat{y} = \text{sgn}(\langle \mathbf{x}, \mathbf{w} \rangle + b)$$

The parameters  $\mathbf{w}$  and  $b$  will uniquely determine the linear classifier.

**Example 1.2.6** (geometric interpretation). We can interpret  $\hat{y} > 0$  as a halfspace (see CO 250). Then, we can draw something like:



### Proposition 1.2.7

The vector  $\mathbf{w}$  is orthogonal to the decision boundary  $H$ .

*Proof.* Let  $\mathbf{x}, \mathbf{x}' \in H$  be vectors on the boundary  $H = \{\mathbf{x} : \langle \mathbf{w}, \mathbf{x} \rangle + b = 0\}$ . Then, we must show  $\mathbf{x}' - \mathbf{x} = \overrightarrow{\mathbf{x}\mathbf{x}'} \perp \mathbf{w}$ .

We can calculate  $\langle \mathbf{w}, \mathbf{x}' - \mathbf{x} \rangle = \langle \mathbf{w}, \mathbf{x} \rangle - \langle \mathbf{w}, \mathbf{x}' \rangle = -b - (-b) = 0$ . □

Originally, the inventor of the perceptron thought it could do anything. He was (obviously) wrong.

---

### Algorithm 1 Training Perceptron

---

**Require:** Dataset  $(\mathbf{x}_i, y_i) \in \mathbb{R}^d \times \{\pm 1\}$ , initialization  $\mathbf{w}_0 \in \mathbb{R}^d$ ,  $b_0 \in \mathbb{R}$ .

**Ensure:**  $\mathbf{w}$  and  $b$  for linear classifier  $\text{sgn}(\langle \mathbf{x}, \mathbf{w} \rangle + b)$

```

for  $t = 1, 2, \dots$  do
    receive index  $I_t \in \{1, \dots, n\}$ 
    if  $y_{I_t}(\langle \mathbf{x}_{I_t}, \mathbf{w} \rangle + b) \leq 0$  then
         $\mathbf{w} \leftarrow \mathbf{w} + y_{I_t} \mathbf{x}_{I_t}$ 
         $b \leftarrow b + y_{I_t}$ 

```

---

In a perceptron, we train by adjusting  $\mathbf{w}$  and  $b$  whenever a training data feature is classified “wrong” (i.e.,  $\text{score}_{\mathbf{w}, b}(\mathbf{x}) := y\hat{y} < 0 \iff$  the signs disagree).

The perceptron solves the feasibility problem

$$\text{Find } \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R} \text{ such that } \forall i, y_i(\langle \mathbf{x}_i, \mathbf{w} \rangle + b) > 0$$

by iterating one-by-one. It will converge “faster” (with fewer  $t$ -iterations) if the data is “easy”.

Consider what happens when there is a “wrong” classification. Let  $w_{k+1} = w_k + yx$  and  $b_{k+1} = b_k + y$ .

Then, the updated score is:

$$\begin{aligned}\text{score}_{\mathbf{w}_{k+1}, b_{k+1}}(\mathbf{x}) &= y \cdot (\langle \mathbf{x}, \mathbf{w}_{k+1} \rangle + b_{k+1}) \\ &= y \cdot (\langle \mathbf{x}, \mathbf{w}_k + y\mathbf{x} \rangle + b_k + y) \\ &= y \cdot (\langle \mathbf{x}, \mathbf{w}_k \rangle + b_k) + \langle \mathbf{x}, \mathbf{x} \rangle + 1 \\ &= y \cdot (\langle \mathbf{x}, \mathbf{w}_k \rangle + b_k) + \underbrace{\|\mathbf{x}\|_2^2 + 1}_{\text{always positive}}\end{aligned}$$

which is always an increase over the previous “wrong” score.